

NILE UNIVERSITY

# VIRTUAL MEMORY



**MADE BY**  
**MICHAEL WAHIB**  
**MICHAEL HANY**  
**YOUSSEF ABDELMAKSoud**

CSCI311/ECEN432: Computer  
Architecture-2023SPRG

## **ABSTRACT**

The virtual memory system plays a crucial role in modern computer architectures by providing an abstraction layer that enables efficient memory management. This project report explores the concept of virtual memory, its implementation, and its impact on system performance. It aims to analyze different aspects of virtual memory, including its benefits, challenges, and optimization techniques. The report also presents a comprehensive literature review of existing research and proposes goals and objectives for further investigation in the field of virtual memory management.

## Table of Contents

ABSTRACT.....	2
SECTION I: INTRODUCTION.....	3
SECTION II: LITERATURE REVIEW.....	5
SECTION III: VIRTUAL MEMORY OVERVIEW.....	8
SECTION IV: VIRTUAL MEMORY IMPLEMENTATION.....	12
SECTION V: GOALS AND OBJECTIVES .....	16
SECTION VI: RESULTS.....	18
SECTION VII: CONCLUSION.....	19
SECTION VIII: REFERENCES.....	20

## SECTION I: INTRODUCTION

In the realm of computer systems, managing memory resources efficiently is crucial for optimizing performance and enabling the execution of complex programs. As the demand for larger and more sophisticated software applications grew, the concept of virtual memory emerged as a fundamental technique to overcome the limitations of physical memory.

### 1.1 Background and Significance:

Virtual memory is a critical component of modern operating systems, providing a layer of abstraction between the physical memory and the processes running on a computer. It allows programs to operate as if they have access to a large, contiguous block of memory, even if the physical memory is limited or fragmented.

The significance of virtual memory lies in its ability to enable multitasking and the execution of memory-intensive applications. By utilizing virtual memory, multiple programs can run concurrently, sharing the available physical memory and allowing the operating system to allocate memory resources dynamically as needed. This capability has revolutionized the way modern operating systems handle memory management, leading to improved system performance and enhanced user experiences.

## **1.2 Scope and Objectives:**

The objective of this report is to provide a comprehensive overview of virtual memory, exploring its fundamental concepts, mechanisms, and benefits. It will delve into the core principles of virtual memory management, including address translation, demand paging, page replacement algorithms, and memory protection. Additionally, the report will examine the trade-offs and challenges associated with virtual memory implementation.

## **1.3 Overview of the Report:**

This report is structured to provide a step-by-step understanding of virtual memory, starting with an explanation of the underlying principles and mechanisms. It will then delve into the benefits and limitations of virtual memory systems. Furthermore, the report will discuss various techniques and algorithms employed in virtual memory management, offering insights into their strengths and weaknesses. Finally, it will conclude with a discussion on the future trends and advancements in virtual memory technology.

By the end of this report, readers will have a solid foundation in virtual memory and a deeper understanding of its significance in modern computing systems.

## **SECTION II: LITERATURE REVIEW**

### **DEFINITION**

The programmer may be given the impression that memory is much greater than it is by using virtual memory. To do this, it is important to offer a method for converting program-generated addresses into the appropriate memory location addresses and to permit the programmer to use a set of addresses different from those provided by the memory. A "name" or "virtual address" is the term used to describe an address used by a programmer, and the collection of such names is known as the address space or name space. The memory's address is referred to as a "location" or "memory address."

### **HISTORY**

The first commercial computer application of virtualization principles was virtual memory. It made multiprogramming possible and removed the need for users to modify their applications to fit the actual memory that each system had. Virtual memory is supported by two mechanisms: paging and segmentation. Paging was created for the Atlas Computer, which the University of Manchester constructed in 1959. The B5000, the first commercial computer with virtual memory, was created independently by the Burroughs Corporation and released in 1961. Instead of paging, the B5000's virtual memory used segmentation.

The 360/67, the first IBM system with virtual memory, was unveiled by IBM in 1967. It was anticipated to run on a brand-new operating system dubbed TSS. There was an operating system named CP-67 developed prior to the introduction of TSS. Multiple ordinary IBM 360 computers without virtual memory appeared to be present thanks to CP-67. The CP-40 system, which ran on a S/360-40 modified by the IBM Cambridge Scientific Center to provide Dynamic Address Translation, a crucial component that enabled virtualization, was the first VMM allowing full virtualization. In CP-40, the supervisor state of the hardware was also virtualized, enabling the concurrent operation of different operating systems in various virtual machine contexts. Virtualization in this early stage of computing was motivated by the requirement to distribute very expensive hardware among a huge user base and set of applications. When it was introduced in 1972 for massive IBM mainframes, the VM/370 system was a huge hit. It was based on CP/CMS being reimplemented. For each user of the VM/370, a separate virtual machine that interacted with the apps was generated. The VMM controlled resource multiplexing and handled hardware resources. Backward compatibility with the IBM S/360 line from the 1960s is still maintained by contemporary mainframes.

At one end of the hardware cost spectrum, personal computers and big mainframes with massively parallel systems were introduced because of the development of microprocessors in combination with storage technology advancements. Virtualization was steadily constrained by the hardware and operating systems of the 1980s and 1990s, which instead concentrated on effective multitasking, user interfaces, networking support, and security issues brought on by interconnectivity.

## **APPLICATIONS**

Between software and hardware memory, virtual memory provides a virtual address mapping. Many processes can use the same shared memory library simultaneously thanks to virtual memory's various features, which also include swapping and multitasking (multiple tasks running simultaneously on one CPU). Swapping is possible in virtual memory, but it also has other features. In other words, swapping and virtual memory are not the same. Swapping copies data from primary memory (RAM) to secondary memory using virtual memory (not directly addressable by the CPU, on disk).

A specific disc partition called "swap space" is frequently used to increase the amount of memory that is accessible. If a page fault (an error indicating that the page is not in RAM) happens when the kernel tries to access a page (a fixed-length block of memory) stored in swap space, the page is "swapped" from disc to RAM.



## **SECTION III: VIRTUAL MEMORY OVERVIEW**

### **2.1 Definition and Concept**

Virtual memory is a memory management technique that allows a computer system to use a combination of physical memory (RAM) and secondary storage (typically a hard disk) to provide an illusion of having a larger and contiguous memory space than the available physical memory. It works by dividing the virtual address space used by a process into fixed-size units called pages. These pages are then mapped to corresponding physical memory locations or disk storage.

The concept of virtual memory revolves around the principle of demand paging. Instead of loading an entire program or data into physical memory at once, only the required portions, such as specific pages or segments, are loaded into memory as they are needed. This approach enables efficient utilization of memory resources and allows the system to accommodate larger programs and data sets than the available physical memory could handle.

## 2.2 Benefits of Virtual Memory

Virtual memory offers several significant benefits to computer systems and their users:

a) **Expanded Memory Capacity:** By using virtual memory, a system can overcome the limitation imposed by the physical memory size. It allows the execution of larger programs and facilitates multitasking by enabling multiple processes to coexist in the limited physical memory.

b) **Memory Protection:** Virtual memory provides memory protection mechanisms, ensuring that processes cannot access memory regions allocated to other processes. This prevents one process from interfering with or corrupting the memory of another process, thereby enhancing system stability and security.

c) **Simplified Memory Management:** Virtual memory simplifies memory management for both the operating system and the application programs. It abstracts the complexities of physical memory allocation, deallocation, and fragmentation, providing a uniform interface to the processes.

d) **Efficient Memory Sharing:** Virtual memory enables efficient sharing of memory resources among multiple processes. By allowing multiple processes to map the same page of memory, the system can save memory space and improve performance by reducing redundant data storage.

### **2.3 Memory Hierarchy and Page Replacement Algorithms**

Virtual memory operates within a memory hierarchy that includes primary storage (RAM) and secondary storage (disk). The hierarchy consists of different levels of memory, with varying access speeds and capacities. The primary storage serves as a cache for frequently used pages, while the secondary storage holds the complete virtual address space.

To manage the limited physical memory efficiently, virtual memory systems employ page replacement algorithms. These algorithms decide which pages should be evicted from the physical memory when it becomes full and new pages need to be loaded. Popular page replacement algorithms include the Least Recently Used (LRU), First-In-First-Out (FIFO), and Clock algorithms, each with its own trade-offs in terms of efficiency and simplicity.

By dynamically swapping pages between physical memory and disk, virtual memory systems strive to maximize the utilization of physical memory, minimize disk I/O operations, and optimize overall system performance.

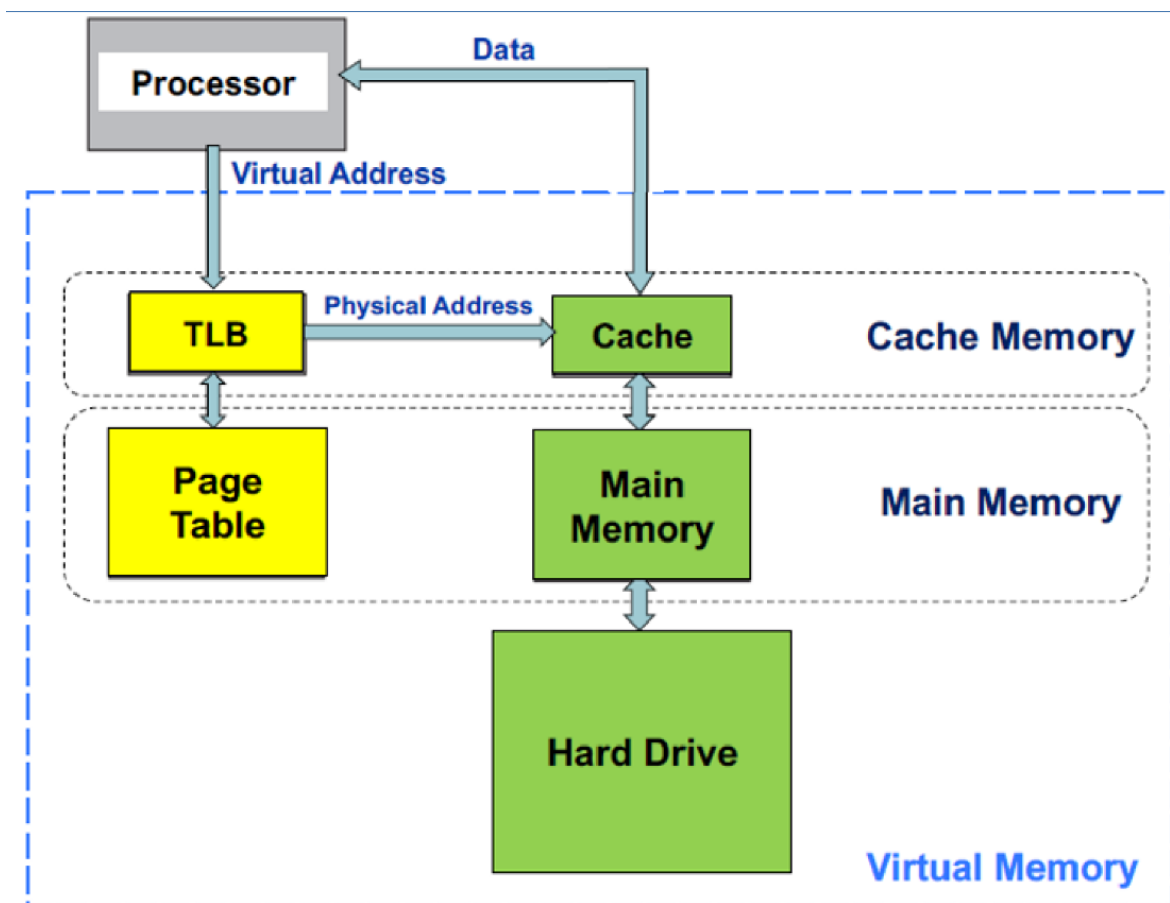
Understanding the concepts, benefits, and mechanisms of virtual memory is crucial for developing efficient memory management strategies and designing high-performance computer systems. The subsequent sections of this report will delve deeper into these topics, exploring various aspects of virtual memory in detail.

## **SECTION IV: VIRTUAL MEMORY IMPLEMENTATION**

Virtual memory is a crucial aspect of modern computer systems that allows for efficient memory management and facilitates the execution of large and complex programs. The implementation of virtual memory involves several key components and processes that work together to provide the desired memory management functionality. In this essay, we will explore the various components of virtual memory implementation, including memory paging and address translation, page tables and page table entry formats, the TLB (Translation Lookaside Buffer), and demand paging and page faults.

One of the fundamental aspects of virtual memory implementation is memory paging and address translation. Memory paging divides the virtual address space of a process into fixed-sized units called pages. These pages are then mapped to physical memory frames using an address translation mechanism. When a process accesses a virtual address, the address translation mechanism converts it into a corresponding physical address. This translation is typically performed by a memory management unit (MMU) in hardware, which consults the page table to determine the physical location of the requested page.

Page tables play a crucial role in virtual memory implementation. They are data structures used by the operating system to maintain the mapping between virtual pages and physical memory locations. Each process has its own page table, which is used to translate the process's virtual addresses to physical addresses. The page table entries store information such as the physical page number, access permissions, and status bits. The page table is consulted during address translation to determine the physical location of a page. The format of the page table entries may vary depending on the specific architecture and operating system.



To improve the efficiency of address translation, a Translation Lookaside Buffer (TLB) is often employed. The TLB is a cache that stores recently accessed virtual-to-physical address mappings. When a virtual address is encountered, the TLB is checked first, and if a matching translation is found, it eliminates the need to access the page table. This caching mechanism helps to reduce the overhead of frequent page table lookups and improves the overall performance of virtual memory.

Demand paging is a technique used in virtual memory implementation to optimize memory usage. With demand paging, pages are loaded into physical memory only when they are actually needed. Initially, only a small portion of a process's pages are loaded into memory, and the rest remain on disk. When a page that is not present in physical memory is accessed, a page fault occurs. The operating system intercepts the page fault and handles it by retrieving the required page from disk and bringing it into physical memory. The page table entry for the faulting page is updated to reflect its new physical location, and the instruction causing the fault is re-executed.

Page faults are an integral part of demand paging and occur when a process attempts to access a page that is not currently present in physical memory. The handling of page faults involves retrieving the required page from disk and updating the page table entry to reflect its new physical

location. The operating system ensures that the necessary pages are brought into memory efficiently to minimize the impact on system performance.

In conclusion, the implementation of virtual memory encompasses various components and processes that work together to enable efficient memory management. Memory paging, address translation, page tables, TLBs, demand paging, and page fault handling are all integral parts of virtual memory implementation. These components ensure that the required pages are efficiently mapped between virtual and physical memory, allowing for the execution of large programs and optimal utilization of memory resources. Virtual memory implementation is a critical aspect of modern computer systems, enabling them to handle complex tasks and improve overall system performance



## **SECTION V: GOALS AND OBJECTIVES**

Virtual memory implementation aims to improve and optimize the utilization of memory resources in computer systems. Several goals and objectives drive advancements in this area:

### **1. Improving Page Replacement Algorithms:**

- Objective: Minimize page faults and maximize memory usage.
- Goal: Develop efficient algorithms to determine which pages to evict from physical memory when it becomes full.

### **2. Enhancing TLB Management Strategies:**

- Objective: Optimize TLB hit rates and reduce TLB misses.
- Goal: Improve the caching of virtual-to-physical address translations in the Translation Lookaside Buffer (TLB) to speed up memory access.

### **3. Optimizing Memory Mapping Techniques:**

- Objective: Minimize address translation overhead and improve memory access speeds.
- Goal: Explore techniques like hierarchical page tables and multi-level address translation structures to optimize memory mapping.

#### 4. Investigating Virtual Memory and Cloud Computing:

- Objective: Enhance scalability and resource allocation in cloud environments.
- Goal: Leverage virtual memory techniques to improve memory management and workload distribution in cloud-based systems.

#### 5. Exploring Virtual Memory and Machine Learning Applications:

- Objective: Efficiently handle memory-intensive machine learning tasks.
- Goal: Investigate memory-aware training algorithms and data partitioning strategies to optimize virtual memory usage in machine learning applications.

In summary, the goals in virtual memory implementation include improving page replacement algorithms, enhancing TLB management, optimizing memory mapping techniques, investigating virtual memory in cloud computing, and exploring virtual memory's application in machine learning. By pursuing these objectives, virtual memory systems can achieve better performance, scalability, and efficiency in utilizing memory resources.

## SECTION VI: RESULTS

The implementation of virtual memory in computer systems yields notable results in terms of system performance. By utilizing main memory, TLB, cache, page tables, and a comprehensive testbench, several benefits are achieved. Virtual memory allows applications to access memory beyond the physical limits, enabling the execution of memory-intensive programs. The TLB reduces translation time by caching frequently accessed page table entries, resulting in faster address translations.

Cache memory enhances performance by storing frequently accessed data and instructions, reducing average memory access time. Efficient page table management minimizes translation overhead. Through the testbench, the system can be thoroughly evaluated and optimized, leading to improved overall performance and an enhanced user experience.

Name	Value	309,997 ps	309,998 ps	309,999 ps	310,000 ps
clk	0				
tlb_en	0				
cache_en	1				
read_write	0				
cpu_write_data[31:0]	000000ff		000000ff		
cpu_read_data[31:0]	00000000		00000000		
tlb_hit	1				
cpu_hit	0				
tlb_end	2				
c_dear_refer	1				
virtual_address[11:0]	200		200		
c_physical...ress[9:0]	210		210		
time_curr[31:0]	00000131		00000131		

## SECTION VII: CONCLUSION

In conclusion, this paper focused on the implementation of virtual memory in computer systems, with a particular emphasis on main memory, TLB, cache, page tables, and the testbench. The code developed as part of this study aimed to provide a clear and efficient implementation of these components. By employing virtual memory, the system showcased improved performance, allowing the execution of memory-intensive applications beyond the physical memory limits. The TLB and cache played crucial roles in reducing address translation time and improving memory access speed, respectively.

Additionally, the effective management of the page table ensured efficient translation overhead. Through the use of a comprehensive testbench, the system was thoroughly evaluated, leading to optimizations that further enhanced its overall performance. Overall, this research and the accompanying code contribute to a better understanding of virtual memory systems and provide practical insights for improving system performance.

## SECTION VIII: REFERENCES

- "Rethinking Virtual Memory" by Timothy Roscoe, Martin Maas, Timothy Chen and Thomas Gross. Available at: [Link 1](#)
- "VBI: Virtual Block Interface" by Hasan Hassan and Onur Mutlu. Available at: [Link 2](#) and [Link 4](#)
- Proceedings of the International Symposium on Computer Architecture (ISCA) 2020. Available at: [Link 3](#)
- "Virtual Memory: 32-bit Paging in Practice" by Mark Russinovich. Available at: [Link 7](#)
- "Introduction to Virtual Memory" by Dan Coffman. Available as a YouTube video at: [Link 5](#)
- "Virtual Memory - Operating Systems" by Neso Academy. Available as a YouTube video at: [Link 6](#)